What IronClaude is

IronClaude is a set of lessons learned about using AI well, wrapped into skills, so that long runs give reliable, repeatable output.

> **Skill:** a packaged set of instructions Claude loads for one kind of work, such as planning, debugging or review. You start most of them with a slash command, for example `/sc:workflow`.

The routing guide, `HOW-WE-WORK.md`, tells you which skill or command to use for each kind of work, and in what order. `superclaude install` copies it to `~/.claude/HOW-WE-WORK.md`. Keep it open while you work: this deck explains the *why*, and the guide holds the exact commands and flags.

> *As it was put in the onboarding meeting: think of the skills as the basic instructions an AI needs to give reliable, repeatable output over a long run, so you can set it going and trust what comes back.*

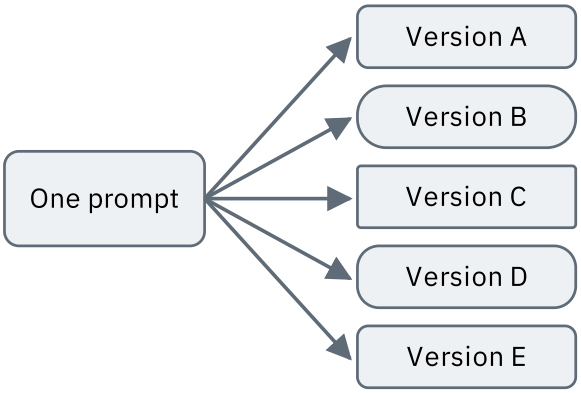
Where anything you hear or read differs from the guide, the guide is the authority.

The problem: consistency

Ask an agent to do the same thing five times and you get five different versions. Each one roughly does what you asked, in a slightly different way, and each one skips something different.

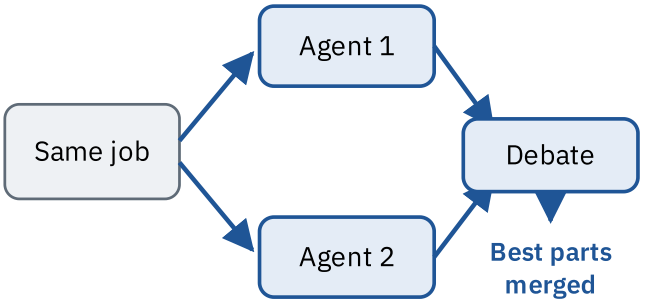
> *In the onboarding meeting the view was that, only a few model generations ago, the best models followed instructions about 70% of the time. That is an opinion from the meeting, not a measured figure, but it explains the design: the heavier skills were built for it, so they add gates, run important work twice, and make agents debate the results.*

Asked once, five times



Five different answers; each misses something different.

Two agents, then a debate



The compare step catches what one run missed.

D1. One prompt fans out to five different results; two agents plus a debate keep the best parts.

How we work as a team

Tokens are cheap; side quests are expensive

> *The advice in the onboarding meeting: do not get pulled into optimizing the harness. Buying more capacity costs less than the downtime of a rabbit hole.*

Drive your work through the commands

Pick the level and the command from the guide before you start (slide 5).

> *In the meeting, about 90% of recent work was said to run through these commands, and the aim was stated plainly: if the commands keep proving to be overkill for your work, the skills should change to fit it.*

Important work gets checked twice

The heavier skills have two agents do the same job, often on different models, and then compare or debate the results. `/sc:adversarial` runs that debate; `/sc:reflect` checks finished work with several independent reviewers.

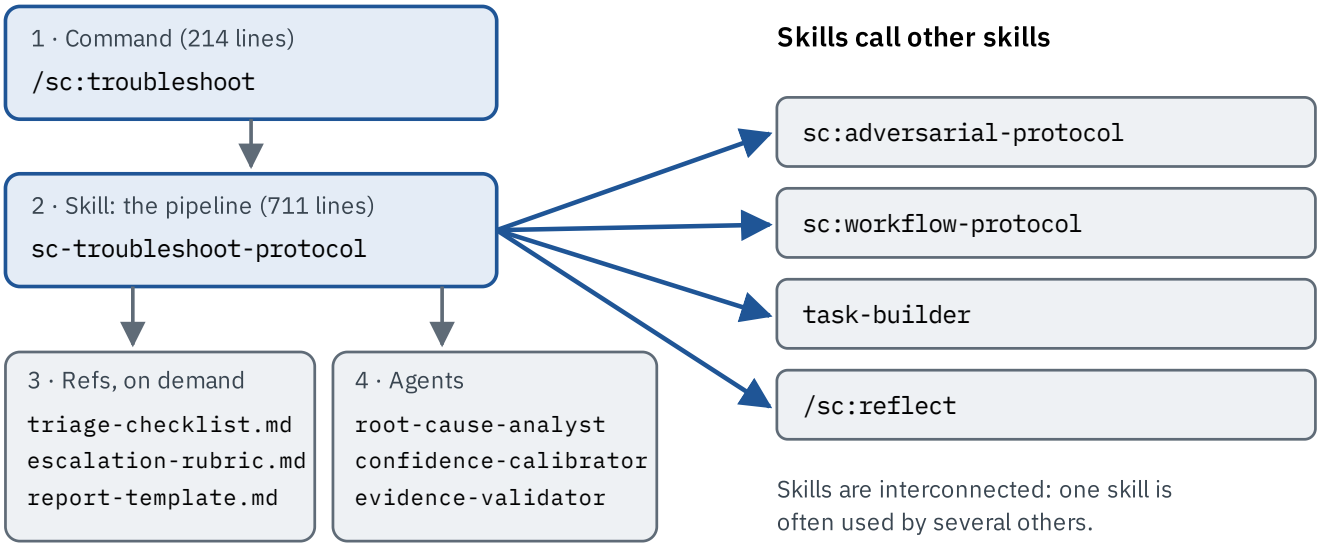
The guide's rules

- Write it down: plans, research and findings are saved to files.
- Let a skill finish; do not repeat its own checks.
- Read a plan before you run it.
- Do only what was asked; suggest anything else.

Anatomy of a command

A command is a thin front door. It tells you and the agent how to use the skill and where the skill lives. The skill holds the actual pipeline.

> **Ref:** an extra instruction file inside a skill's `refs/` folder. **Progressive disclosure:** the skill loads each ref only at the step that needs it, so instructions you do not need never take up room in the conversation.



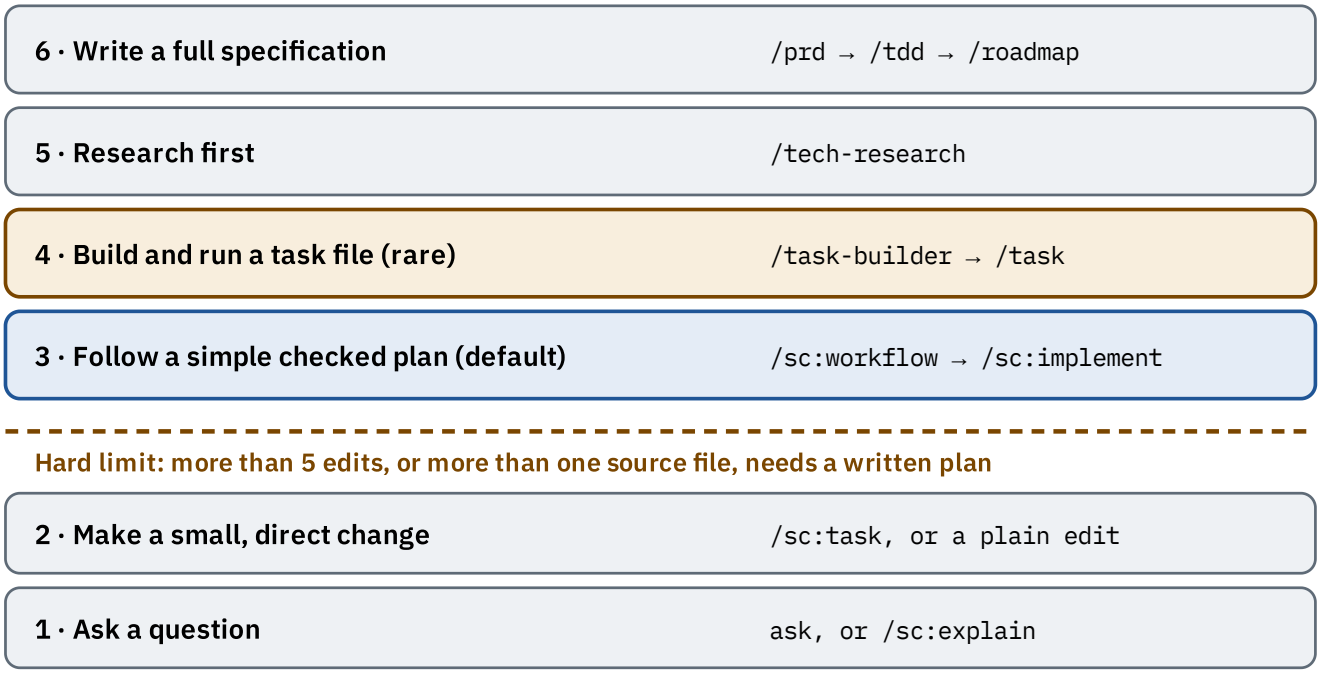
D2. Command, skill, refs (templates such as `report-template.md` are ref files) and agents, with the skills `/sc:troubleshoot` calls.

> *From the meeting: the skills are interconnected, so before you change one, know that several others probably use it.*

Pick your level

Every job falls into one of six levels. Find the first one that describes your job. If you are unsure between levels 3 and 4, choose level 3. If you are unsure between any other two, choose the heavier one.

Hard limit: if the work needs more than 5 edits or touches more than one source file (its matching test file does not count), it must go through a written plan: level 3 by default, level 4 only for the large spec-driven builds described in level 4. The one exception is `/sc:cleanup` and `/sc:improve` run on a folder (guide 3.7): commit first and review `git diff` afterwards.

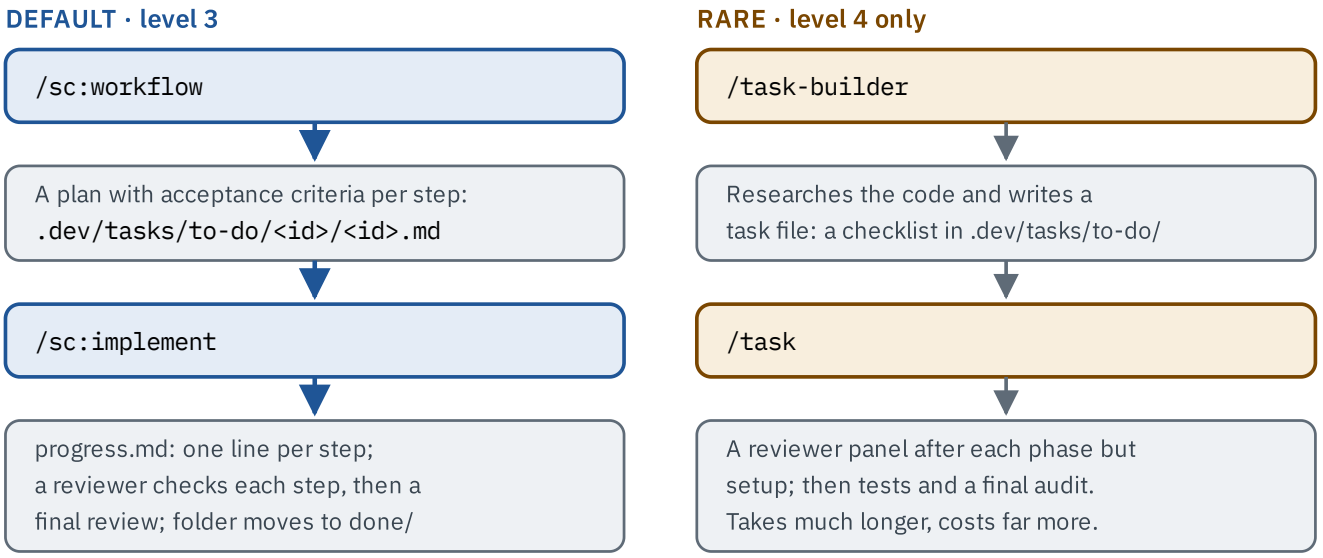


D3. The six levels. Above the hard limit, level 3 is the default; level 4 is only for the large spec-driven builds described in level 4.

Level 4 is for building a whole project or major component from scratch, from a spec complex enough that it nearly warrants a PRD and TDD, when you or your lead chose the extra time and tokens. Level 3 is the default for any planned change, including work that touches many files or takes more than one session.

Two planners

`/sc:workflow` with `/sc:implement` is the default. `/task-builder` with `/task` is the rare, heavy option for level 4.



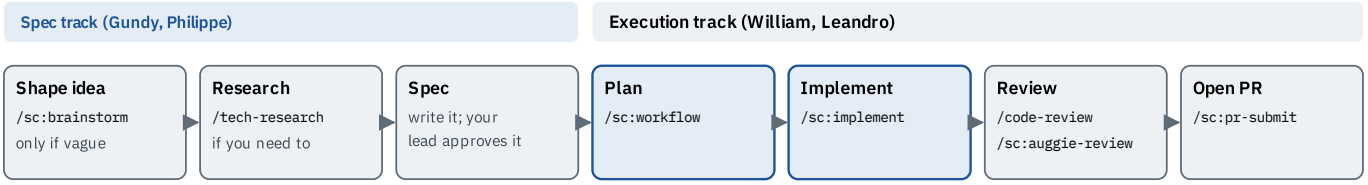
D4. What each pair leaves on disk.

- **Run each plan with its own partner.** Folders named `TASK-WF-*` are `/sc:workflow` plans: run them with `/sc:implement`, never with `/task`. A task file from `/task-builder` runs with `/task`.
- Agents must never choose `/task-builder` or `/task` as the planner on their own, unless a skill or command you started runs them as one of its own steps (slide 13).

> *In the meeting `/sc:workflow` was described as "task builder light": faster and cheaper. It assumes you have done your homework, or lets you tell it plainly in a prompt what to plan. The rule of thumb given there: a plan made by an `sc:` command is followed up by an `sc:` command.*

From idea to merged PR

Most features follow the same path. Skip the steps you do not need: brainstorm only if the idea is vague, research only if you do not know how the current code works.



A new product, or a major feature other people build against (level 6): /prd, then /tdd, then /roadmap, then /task per phase (guide 3.1).
A level-4-sized build swaps Plan and Implement for /task-builder and /task (guide 3.13).

D5. The guide's feature path (3.2), with the two tracks from the onboarding split overlaid.

Spec track (Gundy, Philippe)

Your work starts with a goal that still has to become a spec. These are the tools for that.

Shape a vague idea

`/sc:brainstorm` asks you questions, has several agents propose approaches, and merges them into requirements. It stops there; you decide the next step.

Find out what exists

`/tech-research` investigates *our* codebase. `/sc:research` searches the web. They look alike; they are not.

Check it is not already built

`/sc:recommend "<goal>" --plugin` searches the Claude plugin marketplaces and community skill repositories instead of the local commands.

> *The meeting's advice: before you build anything, look for someone who already built it, then fork it or use it.*

Get an expert critique

`/sc:spec-panel` has a panel of named requirements experts critique a spec. It only reports; you decide what to apply.

When are `/prd` and `/tdd` warranted? Only for level 6: a new product, a new component, or a major feature other people will build against. For a feature on an existing system, write a short spec and plan it with `/sc:workflow` (guide 3.2).

Execution track (William, Leandro)

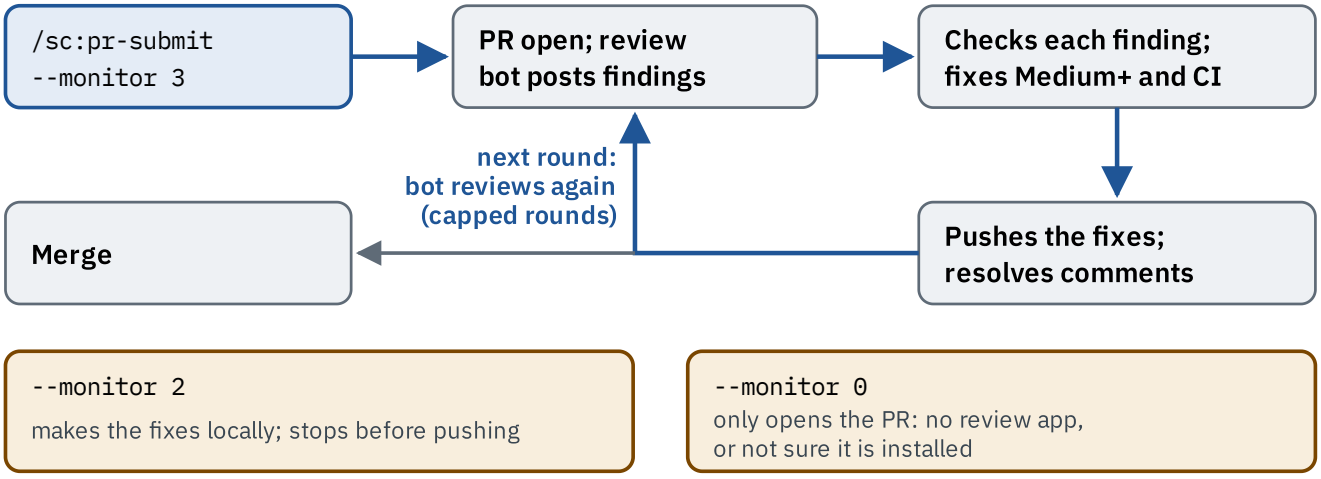
- **Read the plan first.** Check that each step's acceptance criteria clearly say what "done" means.
- **Run it:** `/sc:implement .dev/tasks/to-do/<id>/<id>.md`. It builds one step at a time and a reviewer checks each step against its criteria.
- **The retry ladder:** when a step fails its check, it gets one fix pass, then up to three new attempts written from stronger models' advice. The step's acceptance criteria never change.
- **If interrupted,** run the same command again; it continues from the first unfinished step.

STOP CODE	WHAT IT MEANS	WHAT YOU DO
`E-HALT-QA`	A step still fails its check after the retries.	Fix the code and resume, or have your lead explicitly accept the explained gap.
`E-RESUME-NEEDED`	Something outside the step failed, such as the reviewer crashing.	Run the same command again.
`E-PLAN-CHANGED`	The plan was edited after the first step started.	Do not edit plans once a run has started.

> *Why plans never change after step one, as the meeting put it: an agent that is allowed to edit its own success criteria will rewrite them so that it always succeeds.*

Review and ship

- **Commit locally** on your feature branch. Do not push yet.
- **Bug review:** `/code-review`. **Security review** if you touched login, permissions, secrets, user input or file paths: `/security-review`. Commit the fixes.
- **Branch review:** `/sc:auggie-review` on the whole branch, kept local with `--no-post-pr`.
- **Push only after your lead approves.**
- **Open the PR** with `/sc:pr-submit --monitor 3`.



D6. The PR loop at the setting the guide says to pass, and its two alternatives.

`--monitor 3` is the setting to pass, and the default. After your lead's approved push, it pushes its own review fixes without asking again. If your lead wants to see those fixes first, use `--monitor 2`. Without the Augment Code review app on the repository, or if you are not sure it is installed (the repository's Settings > GitHub Apps page lists it), use `--monitor 0`, which only opens the PR.

Where things live

Work is written to files

Plans, research and findings are saved to files, most under `.dev/`, not only into the chat. If a session dies, the files are what survives.

`superclaude update` overwrites

It overwrites every file it ships, including an edited `~/.claude/CLAUDE.md`, keeps no backup, and re-copies each shipped skill folder. Copy any file you changed or added before you run it.

Resume after a crash

Before you need it, tag the conversation with `/ccsession-tag <name>`; later, from a terminal in the same folder, run `ccsession <name>`.

This always works: run the same command again on the work's own files, `/sc:implement` on the plan or `/task` on the task file.

Not a resume method

`/sc:save` and `/sc:load` store and load project notes. They do not resume a task or a plan.

Being worked out

Two topics came up in the onboarding meeting that have no written rule yet. Each has its own issue. Until they are settled, ask your lead.

Worktree and branch protocol

> **Worktree:** a second checkout of the same repository in its own folder, on its own branch, so parallel sessions do not touch each other's files.

When a session should get its own worktree and branch, where they live and how they are named is being worked out in issue #289.

Contributing rules

How team members propose or change skills, commands and agents, and what evidence a change needs, is being worked out in issue #290.

Traps

Look-alike commands (the guide explains each difference)

`/task`	`/sc:task`
`/sc:workflow` + `/sc:implement`	`/task-builder` + `/task`
`/sc:design`	`/tdd`
`/roadmap`	`/sc:roadmap`
`/tech-reference`	`/sc:document`
`/tech-research`	`/sc:research`
`/sc:cleanup-audit`	`/sc:cleanup`

An agent picking the heavy planner. The guide's rule: an agent must never choose `/task-builder` or `/task` as the planner on its own. It may use them only when you name the command, or after you confirm once you have been told it takes much longer and costs far more than `/sc:workflow`. This does not apply when a skill or command you started runs them as one of its own steps.

> *The sign, as described in the meeting: a long run starts and many reviewer or QA subagents appear that you never asked for.*

`/sc:save` is not a way to resume. Use the steps on slide 11.

When unsure

Ask for a recommendation

`/sc:recommend "<what you want to do, and why>"` picks the right command and writes a ready-to-paste prompt.

See what is installed

`/sc:help` lists the installed commands and skills and shows the six levels.

Ask your lead

Before starting anything large, or whenever this deck and the guide leave you unsure. If you work alone, read "your lead" as yourself.

Cheat sheet

LEVEL OR JOB	COMMANDS	GUIDE SECTION
1 · Ask a question	Ask in plain English, or `/sc:explain`	1, 3.5
2 · Small, direct change	`/sc:task`, or a plain edit, then run the test	1, 3.3
3 · Checked plan (default)	`/sc:workflow` → `/sc:implement`	1, 3.2, 3.3
4 · Task file (rare)	`/task-builder` → `/task`	1, 3.13
5 · Research first	`/tech-research`, then `/sc:workflow`	1, 3.5
6 · Full specification	`/prd` → `/tdd` → `/roadmap` → `/task`	1, 3.1
Fix a bug	`/sc:troubleshoot`	3.4
Review and ship	`/code-review`, `/sc:auggie-review`, `/sc:pr-submit --monitor 3`	3.11
Resume	The same command again, or `ccsession <name>`	3.12
Unsure	`/sc:recommend`, `/sc:help`, your lead	4

Full guide: HOW-WE-WORK.md on GitHub, installed at `~/.claude/HOW-WE-WORK.md`.

Built from HOW-WE-WORK.md at commit 070399eb. Note: this deck already follows the Part A guide changes (the level wording, the hard limit, section 3.13 and the default planner), which are not yet committed; until they are, the guide on GitHub does not show them.

Source of truth for every command and flag: HOW-WE-WORK.md.